

# CIS 575. Introduction to Algorithm Analysis

## Material for April 19, 2024

### Depth-First Search on Directed Graphs: An Application

©2020 Torben Amtoft

The topic of this note is covered in *Cormen's* Section 20.4.


## 1 Topological Sort

For a **directed** graph  $(V, E)$ , *topological* sort is the problem of listing the nodes in a order that “respects” the directed edges, in the sense that

whenever there is an edge from  $u$  to  $w$  then  $u$  must be listed before  $w$ .

If the nodes represent tasks, and an edge from  $u$  to  $w$  represents that  $u$  is a prerequisite for doing  $w$ , a topological sort will show one possible linear schedule for the tasks.


For example, consider the graph




which allows *two* topological sorts: ABCD, and ACBD.

Obviously, if the graph has a cycle then no topological sort is possible. But we shall now show that for an **acyclic** graph it is *always* possible to construct a topological sort.

To do so, let us recall that the DFS algorithm from a directed graph constructs a tree, together with back edges, forward edges, and cross edges. We looked at the example



but observe that a back edge, going from a node to a tree ancestor, will *always* cause a cycle. Let us therefore remove the back edges, to arrive at



We shall use this example to derive a general recipe for how to find an order that respects each edge. With such an edge going from  $u$  to  $w$ , there are 3 cases to consider:

- For a *tree edge* (such as from E to F), where  $w$  is discovered only when  $u$  examines it, it will still be the case that (due to nodes being processed using a stack discipline) that  $w$  finishes before  $u$  does.
- For a *forward edge* (such as from B to D), even though  $w$  was discovered after  $u$  was discovered,  $w$  has finished before  $u$  gets to examine  $w$ .
- For a *cross edge* (such as from F to B),  $w$  finished even before  $u$  was discovered.

We see that in *all cases*, the node  $w$  was finished *before*  $u$  was finished. Hence, we can get a topological sort by listing the nodes in *decreasing* order of their finish time. Thus, if we print a node whenever it finishes, the output will be a topological sort in *reverse* order (a later processing can reverse it back).

Looking back at the generic DFS algorithm, this shows that the **PostNode** action should be “print the node”. The other non-trivial action is **OtherEdge** which must report an error for a back edge, but ignore a forward edge or a cross edge. Thus the following instantiation of the DFS algorithm (which inherits its running time  $\Theta(|V| + |E|)$ ) implements (reverse)

topological sort:

```
TOPOLOGICALSORT( $u$ )
   $color[u] \leftarrow \text{gray}$ 
  foreach  $(u, w) \in E$ 
    if  $color[w] = \text{white}$ 
      TOPOLOGICALSORT( $w$ )
    else
      if  $color[w] = \text{gray}$ 
        // back edge
        Error: cyclic graph
   $color[u] \leftarrow \text{black}$ 
  print  $u$ 
```

For our example, the algorithm will output the list DCBFHGEA and the reverse list

AEGHFBCD

is indeed a topological sort (several others exist).