

Flow Networks

Setting: a **directed** graph where

- ▶ each edge e has positive **capacity** $C(e)$
- ▶ one node is **source** (often called s)
- ▶ another node is **sink** (often called t)

A **flow** assigns to each edge e a number $F(e)$ with

$$0 \leq F(e) \leq C(e)$$

which satisfies **flow conservation**:

*for all nodes except source and sink,
sum of **incoming** flow **equals** sum of **outgoing** flow*

Value of flow: sum of **outgoing** flows from **source**

- ▶ minus sum of incoming flow, but typically zero
- ▶ will equal sum of incoming flow to sink

Goal: find flow with **maximum** value

Finding Maximal Flow

An **augmenting path** is an

acyclic path from source to sink

Naive attempt: if **no** augmenting path then **zero** flow; otherwise

1. find augmenting path π
2. let m be **minimum** capacity of edges in π
3. for each edge in π , **subtract** m from its capacity (remove edge if capacity is now 0)
4. **recursively** build flow F for resulting network (with smaller total capacity so recursion will terminate)
5. return F , after adding m to all edges in π

This approach

- ▶ often works
- ▶ but often does **not** work

Ford-Fulkerson Method

To allow us to (partially) **undo** flow assignments, we augment step 3:

- ▶ for each edge in π , **subtract** m from its capacity (remove edge if capacity is now 0)

so that for each edge in π it **also**

- ▶ adds an edge with capacity m in the **reverse** direction (or adds m to an already existing edge)

This is the **Ford-Fulkerson** method which

- ▶ is highly **non**-deterministic
- ▶ returns a **maximum** flow **when** (if!) it terminates
- ▶ **always** terminates when capacities are **integers**
- ▶ never decreases flow of edge from source or into sink

Running time for **integer** network:

- ▶ to find augmenting path: $O(|E|)$ (may use DFS)
- ▶ number of iterations: $O(M)$ with M the max flow

Thus the total running time is in **$O(M|E|)$**

The Edmonds-Karp Algorithm

Recall **Ford-Fulkerson**: we repeatedly

1. select arbitrary **augmenting path**
2. update network so as to allow us to (partially) **undo** flow assignments.

With M the **maximum** flow, this will

- ▶ terminate after at most M iterations
- ▶ but M iterations may actually be needed!

To reduce the risk of many iterations,

*choose an augmenting path with **fewest** edges.*

This is the **Edmonds-Karp** algorithm:

- ▶ we then need at most $O(|V| \cdot |E|)$ iterations
- ▶ and inherit the bound of M iterations
- ▶ with each iteration taking time in $O(|E|)$

The **total running time** is thus in

$$O(\min \left\{ \begin{array}{l} |V| \cdot |E|^2 \\ M \cdot |E| \end{array} \right\})$$

Cuts

A **cut** in flow network (V, E) is a subset U of V where

- ▶ the source is in U
- ▶ but the sink is **not** in U .

The **capacity** of a cut U is given by

$$C(U) = \sum_{u \in U, w \notin U} C(u \rightarrow w)$$

A cut with **smallest** capacity is a “**bottleneck**”; indeed:

- ▶ with M the maximum (value of a) flow
- ▶ and C the minimum (capacity of a) cut

we have $M = C$ as we shall now argue.

Flows Cannot Exceed Cuts

$$\begin{aligned} V(F) &= \sum_{v \in V} F(s, v) - \sum_{v \in V} F(v, s) + 0 \\ &= \sum_{v \in V} (F(s, v) - F(v, s)) + \\ &\quad \sum_{u \in U \setminus \{s\}, v \in V} (F(u, v) - F(v, u)) \\ &= \sum_{u \in U, v \notin U} F(u, v) + \sum_{u \in U, v \in U} F(u, v) \\ &\quad - \sum_{u \in U, v \in U} F(v, u) - \sum_{u \in U, v \notin U} F(v, u) \\ &= \sum_{u \in U, v \notin U} F(u, v) - \sum_{u \in U, v \notin U} F(v, u) \\ &\leq \sum_{u \in U, v \notin U} F(u, v) \leq \sum_{u \in U, v \notin U} C(u, v) = C(U) \end{aligned}$$

Finding Bottlenecks

Theorem: with M the maximal flow ,

there exists a cut U with capacity M .

Proof for $M = 0$: then there is no augmenting path, so

$$U = \{v \mid \text{path from source to } v\}$$

- ▶ is a **cut** since it contains the source but **not** the sink
- ▶ has **capacity zero** since there is **no** edge from a node $\in U$ to a node $\notin U$.

Proof for $M > 0$: apply Ford-Fulkerson until **no** augmenting path; there is (cf above) a **cut** U that

- ▶ has capacity zero in the resulting network
- ▶ but then U has **capacity M** in the original network.

We see that we can find a minimum cut **without** (exponential) brute-force search.

Bipartite Matching

A **Bipartite Graph** is an **undirected** graph (V, E) where there exists V_1, V_2 such that

- ▶ $V_1 \cup V_2 = V$ and $V_1 \cap V_2 = \emptyset$
- ▶ each edge $\in E$ is:
 between a node $\in V_1$ and a node $\in V_2$

A **Matching** is a subset E' of E such that for all $v \in V$:
 *v is the end point of **at most one edge** in E'*

Goal: find a matching that is **maximal**
in that no other matching has more edges

Construct Maximal Matching

From **bipartite graph** $(V_1 \cup V_2, E)$ we
construct a **flow network** (V', E') :

- ▶ $V' = V_1 \cup V_2 \cup \{s, t\}$ where s and t are **fresh**
- ▶ s is the source; t is the sink
- ▶ each edge has **capacity 1**
- ▶ E' is given by

$$\begin{aligned} & \{(s \rightarrow v_1) \mid v_1 \in V_1\} \\ & \cup \{(v_1 \rightarrow v_2) \mid (v_1, v_2) \in E\} \\ & \cup \{(v_2 \rightarrow t) \mid v_2 \in V_2\} \end{aligned}$$

Fact: the network

has a **flow** with value M

iff the bipartite graph has a **matching** with M edges.

Thus a maximal matching can be

- ▶ constructed by the Ford-Fulkerson method
- ▶ where a node once matched will remain matched (though may switch partner)